

# Automatic Generation of Matrix-Free Routines for PDE Solvers with Devito via PETSc

Z. Leibowitz<sup>1</sup>, R. Nelson<sup>1</sup>, F. Luporini<sup>2</sup>, M. Louboutin<sup>2</sup>,  
M. Knepley<sup>3</sup>, L. Mitchell, J. Betteridge<sup>2</sup>, E. Caunt<sup>2</sup>, G.  
Bisbas<sup>1</sup>, M. Piggott<sup>1</sup>, G. Gorman<sup>1,2</sup>

<sup>1</sup> Imperial College London

<sup>2</sup> Devito Codes

<sup>3</sup> University at Buffalo

PETSc User Meeting, May 20-21, 2025

## Talk Outline

- Why do people care about DSLs?
- What is Devito?
- What are some things Devito is currently lacking?
- What are my new contributions to the software? What are my updates since last year?
- Results
- Next steps

## Traditional approach to solving PDEs

$$m \frac{\partial^2 u}{\partial t^2} + \eta \frac{\partial u}{\partial t} - \Delta u = 0$$



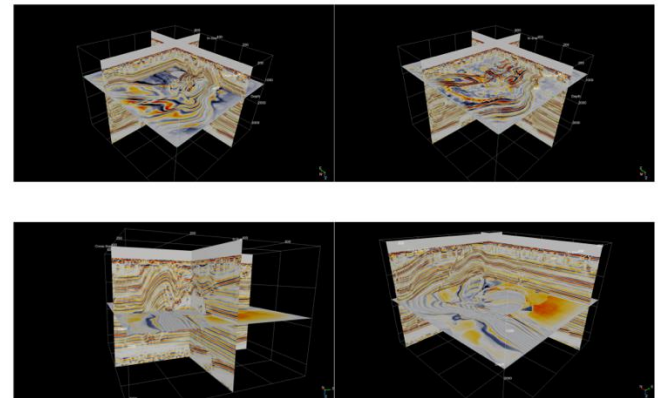
```
void kernel(...) {  
    ...  
    <impenetrable code with aggressive performance optimizations  
    written by rockstars, gurus, ninjas, unicorns and celestial  
    beings>    ...  
}
```

## Domain Specific Languages (DSLs) Manifesto

- Decoupling mathematics from HPC
- Productivity, portability, performance
- Write once, run everywhere
- Do in  $O(\text{days})$  what would normally take  $O(\text{months})$

## Introduction to Devito

- A DSL and compiler to express explicit finite difference operators.
- Embedded in Python – easy to learn, symbolic API.
- Generates optimised C code for multiple architectures:
  - CPUs
    - OpenMP + MPI
  - GPUs
    - OpenMP/OpenACC + MPI
    - CUDA/HIP/SYCL (proprietary)
- Many users in both academia and industry.
- Open source platform – MIT license.



Louboutin et al.: Synthetic 3D anisotropic subsurface model simulating a realistic industry-scale TTI problem

## The Devito Vision



$$m \frac{\partial^2 u}{\partial t^2}(x, t) + \eta \frac{\partial u}{\partial t}(x, t) - \nabla^2 u(x, t) = 0$$

*User writes ...*

```
eqn = m * u.dt2 + eta * u.dt - u.laplace
```

*Devito creates ...*

```
void kernel (...) ...
```

## Example – simple 2D diffusion operator

## Example – simple 2D diffusion operator

**Step 1** - Import Devito:

```
from devito import Grid, TimeFunction, Eq, solve, Operator
```



## Example – simple 2D diffusion operator

**Step 1** - Import Devito:

```
from devito import Grid, TimeFunction, Eq, solve, Operator
```

**Step 2** - Create computational domain:

```
grid = Grid(shape=(11, 11), extent=(1., 1.))
```

## Example – simple 2D diffusion operator

**Step 1** - Import Devito:

```
from devito import Grid, TimeFunction, Eq, solve, Operator
```

**Step 2** - Create computational domain:

```
grid = Grid(shape=(11, 11), extent=(1., 1.))
```

**Step 3** - Define the functions present in the PDE:

```
f = TimeFunction(name='f', grid=grid, space_order=so, ...)
```

## Example – simple 2D diffusion operator

**Step 1** - Import Devito:

```
from devito import Grid, TimeFunction, Eq, solve, Operator
```

**Step 2** - Create computational domain:

```
grid = Grid(shape=(11, 11), extent=(1., 1.))
```

**Step 3** - Define the functions present in the PDE:

```
f = TimeFunction(name='f', grid=grid, space_order=so, ...)
```

**Step 4** - Define the PDE, explicit FD scheme and boundary conditions:

```
eqn = Eq(f.dt, 0.5 * f.laplace)  
update = Eq(f.forward, solve(eqn, f.forward))  
bcs = ... # BCs expressed in symbolic form
```

## Example – simple 2D diffusion operator

**Step 5** - Generate C kernel through a sequence of compilation passes:

```
op = Operator(update, bcs, ...)
```

**This will generate OpenMP + MPI code!**

## Example – simple 2D diffusion operator

**Step 5** - Generate C kernel through a sequence of compilation passes:

```
op = Operator(update, bcs, ...)
```

**This will generate OpenMP + MPI code!**

**Step 6** - Just-in-time (JIT) compile and execute:

```
op(t=timesteps, dt=dt)
```

## Example – simple 2D diffusion operator

**Step 5** - Generate C kernel through a sequence of compilation passes:

```
op = Operator(update, bcs, ...)
```

**This will generate OpenMP + MPI code!**

**Step 6** - Just-in-time (JIT) compile and execute:

```
op(t=timesteps, dt=dt)
```

Low level C loop structure automatically generated:

```
for (int time = time_m, t0 = (time)%2, t1 = (time + 1)%2; time <= time_M; time += 1, t0 = (time)%2, t1 = (time + 1)%2)
{
    for (int x = x_m; x <= x_M; x += 1)
    {
        for (int y = y_m; y <= y_M; y += 1)
        {
            f[t1][x + 2][y + 2] = dt*(r2*f[t0][x + 2][y + 2] + (-1.0F)*(r0*f[t0][x + 2][y + 2] + r1*f[t0][x + 2][y + 2]) + ... );
        }
    }
}
```

## Flexibility in terms of the numerical scheme

`f = TimeFunction(name='f', grid=grid, space_order=so, ...)`

**so = 2**

```
for (int time = time_m, t0 = (time)%2, t1 = (time + 1)%2; ...)
{
    for (int x = x_m; x <= x_M; x += 1)
    {
        for (int y = y_m; y <= y_M; y += 1)
        {
            f[t1][x + 2][y + 2] = dt*(r2*f[t0][x + 2][y + 2] +
            (-1.0F)*(r0*f[t0][x + 2][y + 2] + r1*f[t0][x + 2][y + 2]) +
            5.0e-1F*(r0*f[t0][x + 1][y + 2] + r0*f[t0][x + 3][y + 2] +
            r1*f[t0][x + 2][y + 1] + r1*f[t0][x + 2][y + 3]));
        }
    }
}
```

**so = 12**

```
for (int time = time_m, t0 = (time)%2, t1 = (time + 1)%2; ...)
{
    for (int x = x_m; x <= x_M; x += 1)
    {
        for (int y = y_m; y <= y_M; y += 1)
        {
            float r3 = -2.98277780F*f[t0][x + 12][y + 12];
            f[t1][x + 12][y + 12] = dt*(r2*f[t0][x + 12][y + 12] +
            5.0e-1F*(r0*(r3 + (-6.01250601e-5F)*(f[t0][x + 6][y + 12] +
            f[t0][x + 18][y + 12]) + 1.03896104e-3F*(f[t0][x + 7][y + 12] +
            f[t0][x + 17][y + 12]) + (-8.92857143e-3F)*(f[t0][x + 8][y + 12] +
            f[t0][x + 16][y + 12]) + 5.29100529e-2F*(f[t0][x + 9][y + 12] +
            f[t0][x + 15][y + 12]) + (-2.67857143e-1F)*(f[t0][x + 10][y + 12] +
            f[t0][x + 14][y + 12]) + 1.714285710F*(f[t0][x + 11][y + 12] +
            f[t0][x + 13][y + 12])) + r1*(r3 + (-6.01250601e-5F)*(f[t0][x + 12][y + 6] +
            f[t0][x + 12][y + 18]) + 1.03896104e-3F*(f[t0][x + 12][y + 7] +
            f[t0][x + 12][y + 17]) + (-8.92857143e-3F)*(f[t0][x + 12][y + 8] +
            f[t0][x + 12][y + 16]) + 5.29100529e-2F*(f[t0][x + 12][y + 9] +
            f[t0][x + 12][y + 15]) + (-2.67857143e-1F)*(f[t0][x + 12][y + 10] +
            f[t0][x + 12][y + 14]) + 1.714285710F*(f[t0][x + 12][y + 11] +
            f[t0][x + 12][y + 13]))));
        }
    }
}
```

## Limitations to certain application domains



Nicholas Doherty from unsplash

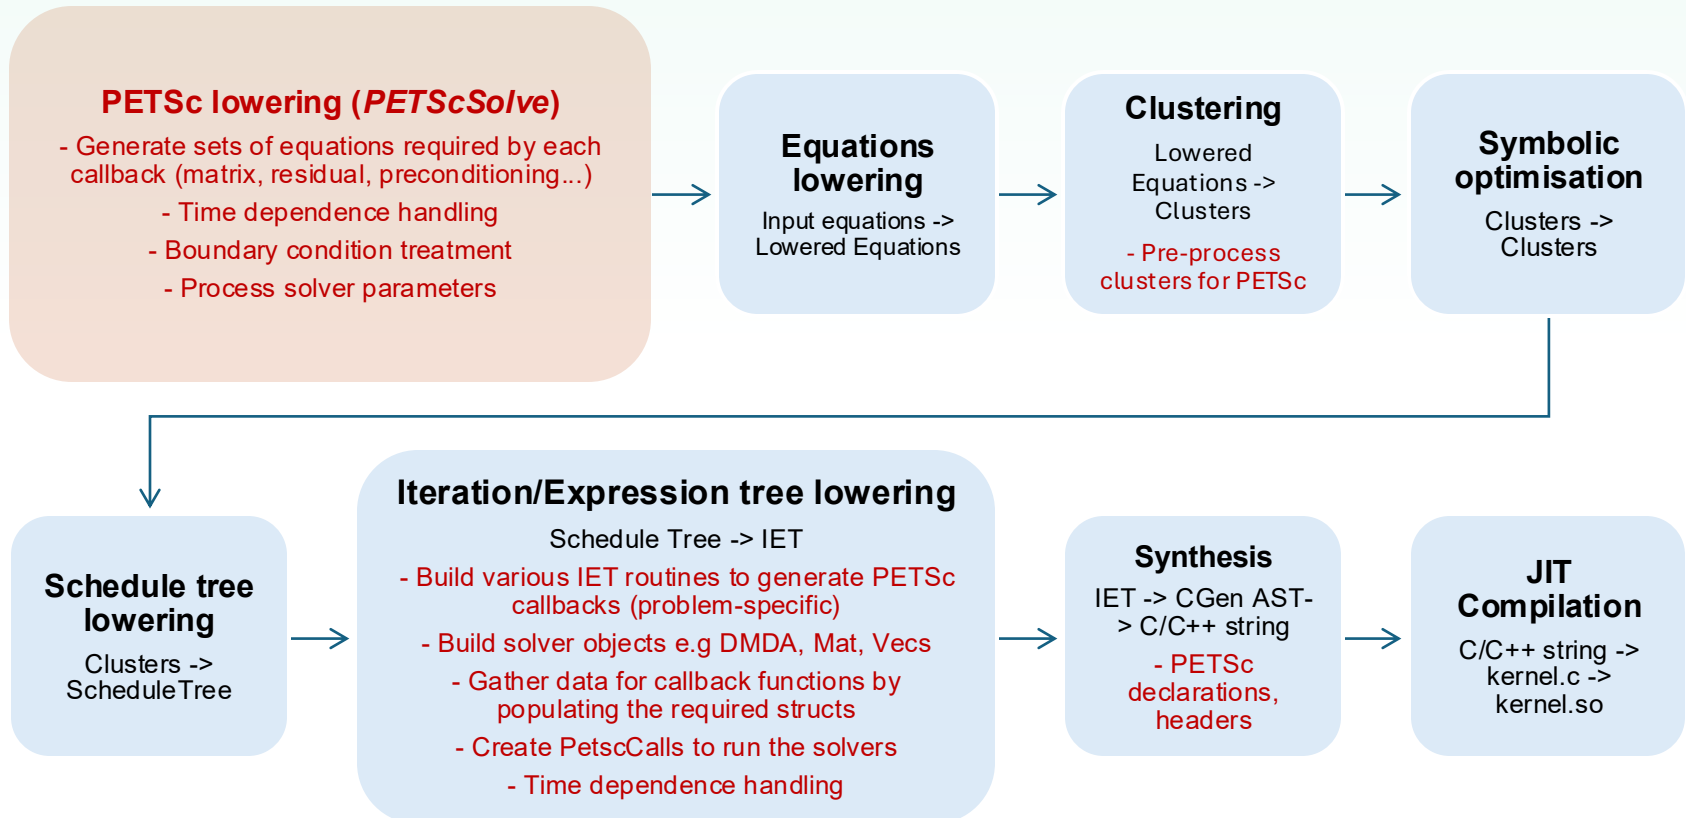
**My work:** Automatically generate matrix-free routines in Devito to interact with PETSc's scalable solvers.



- Devito compiler pass

## Main update since PETSc Meeting 2024

- Automation!!!




## Example: The biharmonic equation as a coupled system

$$\nabla^4 u = f$$

$$u = 0, \quad \nabla^2 u = 0 \quad \text{on } \partial\Omega$$

Substitute  $v = -\nabla^2 u$ , then system has triangular block form:

$$\begin{aligned} -\nabla^2 v &= f \\ -\nabla^2 u &= v \end{aligned} \quad \Leftrightarrow \quad \begin{bmatrix} -\nabla^2 & 0 \\ -I & -\nabla^2 \end{bmatrix} \begin{bmatrix} v \\ u \end{bmatrix} = \begin{bmatrix} f \\ 0 \end{bmatrix}$$



We manufacture an exact solution  $u(x, y) = c(x)c(y)$ , where  $c(x)$  is a 6th-degree polynomial that satisfies the boundary conditions.

## User level code: The biharmonic equation

```
# Build computational domain
grid = Grid(shape=(n,n), extent=(1.,1.))

# Define functions present in PDEs
u = Function(name='u', grid=grid, space_order=2)
v = Function(name='v', grid=grid, space_order=2)
f = Function(name='f', grid=grid, space_order=2)

# Set up RHS
f.data[:] = ...

# Symbolically express the equations
eq_v = Eq(-v.laplace, f, subdomain=grid.interior)
eq_u = Eq(-u.laplace, v, subdomain=grid.interior)

# Employ new API object to trigger the lowering to PETSc
solver = PETScSolve({v: [eq_v], u: [eq_u]},
                    solver_parameters={'ksp_type': 'gmres',
                                       'ksp_rtol': 1e-7,
                                       'ksp_atol': 1e-30, ...})

# Automatically generate the low level C kernel
op = Operator(solver, language='petsc')

# JIT compile and execute the code
op.apply()
```

## Automatically generated PETSc code: The biharmonic equation

```
...
PetscCall(DMDACreate2d(...,2,...,&da));
...
PetscCall(DMSetMatType(da,MATSHELL));
PetscCall(DMCreateMatrix(da,&(J)));
PetscCall(SNESSetJacobian(snes,J,J,MatMFFDComputeJacobian,NULL));
PetscCall(SNESSetType(snes,SNESKSPONLY));
PetscCall(KSPSetType(ksp,KSPGMRES));
PetscCall(MatShellSetOperation(J,MATOP_MULT,...WholeMatMult0));
PetscCall(SNESSetFunction(snes,NULL,WholeFormFunc0,(void*)(da)));
PetscCall(DMSetApplicationContext(da,&(ctx)));
...
PetscCall(DMCreateFieldDecomposition(da,...,&fields,&subdms));
PetscCall(MatShellSetOperation(J,...MatCreateSubMatrices0));
...
PetscCall(SNESSolve(snes,NULL,xglobal));
...
```

```
PetscErrorCode MatCreateSubMatrices0(Mat J, PetscInt nfields, IS *irow, ...)
{
    ...
    PetscCall(PetscCalloc1(nsubmats,submats));
    for (int i = 0; i <= nsubmats - 1; i += 1)
    {
        PetscCall(MatCreate(PETSC_COMM_WORLD,&(block)));
        PetscCall(MatSetSizes(block,...));
        PetscCall(MatSetType(block,MATSHELL));
        ...
        subctx->rows = &(irow[rowidx]);
        subctx->cols = &(icol[colidx]);
        PetscCall(MatSetUp(block));
        submat_arr[i] = block;
    }
    PetscCall(MatShellSetOperation(submat_arr[0],MATOP_MULT,...J00_MatMult0));
    PetscCall(MatShellSetOperation(submat_arr[2],MATOP_MULT,...J10_MatMult0));
    PetscCall(MatShellSetOperation(submat_arr[3],MATOP_MULT,...J11_MatMult0));
    PetscFunctionReturn(0);
}
```


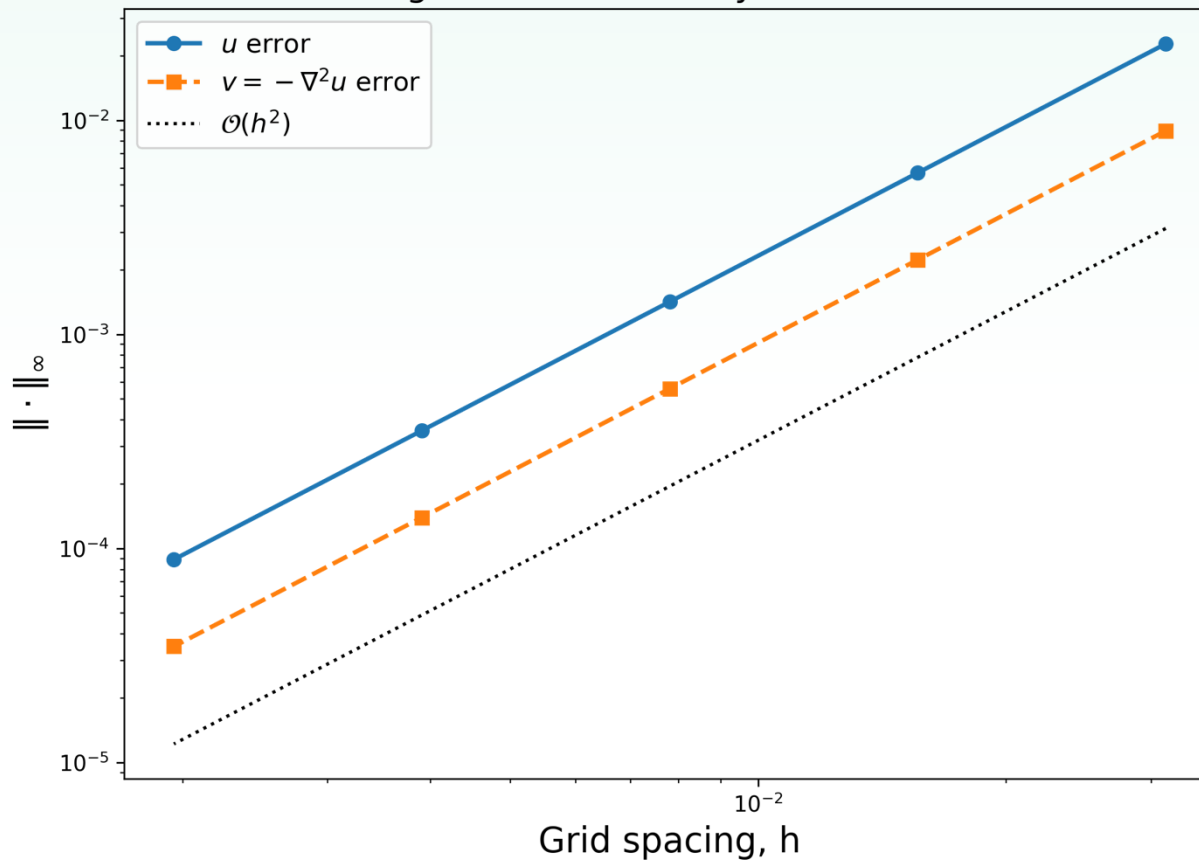
```
PetscErrorCode WholeFormFunc0(SNES snes, Vec X, Vec F, void* ptr)
{
    ...
    for (int ix = ctx0->x_m + ctx0->x_ltkn0; ...)
    {
        for (int iy = ctx0->y_m + ctx0->y_ltkn0; ...)
        {
            f_vu[ix + 2][iy + 2].v = 2.0*(r4*x_vu[ix + 2][iy + 2].v + ... ;
            f_vu[ix + 2][iy + 2].u = 2.0*(r4*x_vu[ix + 2][iy + 2].u + ... ;
        }
    }
    ...
    PetscCall(DMLocalToGlobalBegin(dm,floc,ADD_VALUES,F));
    PetscCall(DMLocalToGlobalEnd(dm,floc,ADD_VALUES,F));
    ...

    PetscFunctionReturn(0);
}
```

```
PetscErrorCode J00_MatMult0(Mat J, Vec X, Vec Y)
{
    ...
    for (int ix = ctx0->x_m + ctx0->x_ltkn0; ...)
    {
        for (int iy = ctx0->y_m + ctx0->y_ltkn0; ...)
        {
            y_v[ix + 2][iy + 2] = 2.0*(r0*x_v[ix + 2][iy + 2] + r1*... ;
        }
    }
    ...
    PetscCall(VecRestoreArray(yloc,&y_v_vec));
    PetscCall(VecRestoreArray(xloc,&x_v_vec));
    PetscCall(DMLocalToGlobalBegin(dm,yloc,ADD_VALUES,Y));
    PetscCall(DMLocalToGlobalEnd(dm,yloc,ADD_VALUES,Y));
    ...
    PetscFunctionReturn(0);
}
```

## Convergence plot: The biharmonic equation

Convergence of the infinity norm for  $u$  and  $v$



## Flexibility

```
solver = PETScSolve(..., solver_parameters={'ksp_type': 'cg', 'pc_type': 'jacobi'})
```

```
solver = PETScSolve([eqn, bc1, bc2, bc3, ...], target=f, ...)
```

```
solver = PETScSolve({f: [eq_f, bc_f1, bc_f2,...], g: [eq_g, bc_g1, bc_g2,...]}, ...)
```

```
solver1 = PETScSolve(eq_u, target=u, ...)  
solver2 = PETScSolve(eq_v, target=v, ...)  
op = Operator([solver1, solver2], language='petsc')
```

And more ...

## Next steps

- Stress testing compiler -> checking for correctness on large range of 2D+3D problems
- Additional solver options: Implement geometric multigrid and other solver and preconditioner options
- Further efficiency and QoL improvements: e.g remove constrained degrees of freedom from global solve with DMDA
- Staggered grids -> DMSTAG?

## Summary

- New compiler passes in Devito now enable iterative solver support via PETSc
- Goal is to produce a flexible, portable, high-level framework
- Enables us to tackle new problem classes, such as incompressible flow in CFD
- Very much in development phase – actively exploring the best ways to interface with PETSc – if you have suggestions or ideas, I'd love to hear them!



GitHub



Website



## Summary

- New compiler passes in Devito now enable iterative solver support via PETSc
- Goal is to produce a flexible, portable, high-level framework
- Enables us to tackle new problem classes, such as incompressible flow in CFD
- Very much in development phase – actively exploring the best ways to interface with PETSc – if you have suggestions or ideas, I'd love to hear them!

**Thank you!**



GitHub



Website